

Distributed System Proofs in Constructor Based Reachability Logic

Paul Pok Hym Ng^a

^aECE, University of Illinois at Urbana-Champaign, ppng2@illinois.edu

Abstract

Proving properties of distributed systems is not an easy task and often using traditional interactive theorem provers is a herculean task. Here we explore proving distributed systems using something known as constructor based reachability logic and rewrite theories. First discovered and designed by Stephen Skeirik [11] for his PhD thesis, reachability logic is a theory generic proof system meaning all tactics and strategies apply for all possible systems (or programming languages). Since then, reachability logic has been further improved and augmented with a tool written in Maude. In addition rewrite theories provide an intuitive description for distributed systems via rewrite rules.

Keywords: Rewrite Theory, Reachability Logic, Maude, Distributed Systems, Interactive Theorem Provers, Theorem Proving, Formal Methods

1 Introduction

Proving distributed systems is complicated due to their non-determinism and non-terminating properties.

There are two approaches to proving distributed systems bounded model checking and theorem proving. The first, bounded model checking, uses languages such as TLA+ to check all possible instances of a distributed system up to a certain bound (number of iterations, number of variables in a system). However this in itself provides a problem as this does not cover all possible states that a distributed system can enter. The latter, theorem proving, solves this issue by performing a mathematical analysis of the distributed system and reasons about these mathematically

expressed properties to prove properties.

Constructor based reachability logic (CBRL) is an instance of the latter (theorem proving) which uses rewrite theories written in the language Maude. In this study we will give a crash course on rewrite theories, CBRL, examples of such proofs with increasing complexity, and finally a comparison to other theorem provers on the market (Isabelle and Coq).

2 Goals

Originally, the goal was to prove the soundness and completeness of a concurrent garbage collection algorithm for actors [9]. In fact, the specification has been completed and will be attached along with all the other specifications and proofs discussed in this paper.

During the early months before Stephen's tool matured, proving, proving such a complex algorithm proved prohibitive. As such, I attempted to rebuild the system in Coq or Isabelle. However, due to my lack of understanding, after several weeks I had yet to prove the simplest Maude module I knew of "Choice" (discussed below).

Therefore as I helped debug, fix, and test `rltool`, the tool that implements CBRL, the goal changed to proving increasingly complex distributed systems. A later section will be dedicated to the troubles and tribulations that have occurred over this journey.

I would like to extend my sincere thanks and gratefulness to Stephen for putting up with all the bugs I discovered and fixing them in a timely manner in order to allow me to complete the proofs shown in this paper.

3 Preliminaries

In this section we will cover a quick rundown of rewrite theories and CBRL. With this knowledge we hope to convey the ideas of CBRL and how to use it. Certain details will be left out for brevity's sake and while important when creating a rewrite theory and using it are not required to understand the ideas behind CBRL and its usage.

3.1 Rewrite Theories

To describe distributed systems we use rewrite theories. A rewrite theory is a 3-tuple (Σ, E, R) . Here Σ represents the function symbols and E represents the equations used to describe the state of a distributed system. In addition, (Σ, E) provides us with all the terms in our initial algebra $T_{\Sigma/E}$. For example, in the theory of Peano numbers this would contain the values $0, s(0), s(s(0)), \dots$. The transitions between states describing the evolution of the distributed system are modeled by R the rewrite rules in our system. Rules can be conditional or unconditional and are applied *modulo* our equations E . This means that a rewrite can apply functions in or use predicates in (Σ, E) to evaluate and trigger the rewrite rule.

3.2 Reachability Logic

CBRL is *theory generic* which allows it to use the same inference rules for different systems with different semantics. This is unlike Hoare logic which requires a redesign of inference rules per language. Often the number of rules are in the high tens and take a long time to prove the completeness/soundness of. In CBRL, as long as we have a suitable \mathcal{R} , rewrite theory, as an input we can use its inference rules.

A reachability logic formula has the following form

$$A \rightarrow^{\circledast} B \iff u \mid \phi \rightarrow^{\circledast} v \mid \psi$$

where A and B are state predicates and the arrow with a star represents a reachability logic formula rather than a rewrite rule.

In the easier case, both A and B have no shared variables. Then given a \mathcal{R} we interpret the terms in A and B in the initial algebra $\mathcal{T}_{\mathcal{R}}$ where a term u can be rewritten to v . Equivalently we can say a state transition from $\mathcal{R} \vdash u \rightarrow v \iff [u] \rightarrow [v]$ (where $[]$ is the equivalence class operator) gives us the result that computation is equivalent to deduction.

Moreover, our terms u and v in a reachability logic formula must be constructors. Where a constructor Ω is a subset of Σ . A simple example is as

follows. In the theory of Peano natural numbers the two constructors are s and 0 the successor. These are not functions on natural numbers but ground terms used to construct a natural number. A term composed completely of constructors is on the left and the right is a term.

$$s(s(s(0))) = s(0) + s(s(0))$$

ϕ and ψ represent a predicate which only use conjunction (\wedge) and disjunction (\vee). These are known as constrained pattern predicates in CBRL and will be explained later.

3.3 Order Sorted Algebras

In Maude all functional modules and modules are described with an order sorted algebra. An order sorted algebra contains the following Σ . $\Sigma = (S, \leq, \Sigma)$ where S is a set of sorts (classes in OOP), \leq a binary relation telling us about the "largeness" of sorts (class inheritance in OOP) and Σ the set of function operators.

Furthermore we must also have the following properties for an algebra (S, Σ) algebra.

1. If $s \leq s'$ then algebra $A' \leq$ algebra A
2. If $f : w \rightarrow s, f : w' \rightarrow s' \in f_{[s]}^{[s_1], \dots, [s_n]}$ (a function f goes from w (w has input sorts s_1, \dots, s_n , and output s has sort s AND a function goes from w' to s' where s' is also a subsort of s), and $\bar{a} \in A^w \cap A^{w'}$ (\bar{a} is a term that exists in both algebras) then we have the result $A_{f:w \rightarrow s}(\bar{a}) = A_{f:w' \rightarrow s'}(\bar{a})$

An example of property 1 is sort *Node* and a sort *Graph*. A natural number is a subsort of integers as the set of all natural numbers is contained within the set of integers. In the case of property 2 we can look at the following example.

$$(5).Nat + (6).Nat = (5).Int + (6).Int$$


Both of these obviously evaluate to the same value, use the same operator, and are different algebras.

This concludes our introduction to order sorted algebras.


3.4 Rewriting Logic

A rewrite theory is used to model a distributed system in our case. \mathcal{R} is a rewrite theory composed of $(\Sigma, E \cup B, R)$ where Σ are function symbols, E are equations, B are axioms, and R are rewrite rules.

All terms in a rewrite theory can be represented as a tree. The tree for $s(0) + (s(0) + s(0))$ is shown below.



All equations or rules that operate on the term will have to work from the leaves of the tree up simplifying each subterm until we have only a root term represented by constructors only. Below we see the steps used by the $+$ operator to simplify the tree above.



Recall that other than R the rest of a rewrite theory is an order sorted algebra. An unconditional rewrite rule is of the form

$$R_i = u \rightarrow v$$

Where a term u can be mutated into term v . An example of this is shown below.

$$N; L \rightarrow L$$

Given a list, with concatenation operator “;”, containing a number N and its tail L (which may be empty), remove the number N from the list. Note that this rule can not apply if the list is empty because we are not able to match an empty list to a non-empty list. In addition to unconditional rules we also have conditional rules which apply only if a condition is true.

$$N; L \rightarrow N; N; L \mid \text{if } N > 0$$

This conditional rule can only duplicate an element in our list if the number N is greater than 0. Furthermore a rewrite theory has the following properties.

1. Unconditional equations $u = v$ must have the same shared variables (no free variables) and no repeated variables.
2. Conditional equations can be re-oriented into conditional rewrite rules
3. The rules are ground coherent with E modulo B

Ground coherence [3] is the property where all ground terms (constructor terms) can be written modulo associativity which is in general undecidable $([t]_A \rightarrow [t']_A)$.

These three conditions give us the initial reachability model $\mathcal{T}_{\mathcal{R}}$ and the initial algebra $T_{\Sigma/E \cup B}$'s

isomorphism to the canonical term algebra $C_{\Sigma/E \cup B}$. A canonical term algebra is the simplest form of representing a term and is most often composed completely by constructors of the algebra.

3.5 Constructor Decomposition

Constructor decomposition is the the subset theory obtained from $(\Sigma, E \cup B)$ which operates completely on the constructors.

More concretely we turn $(\Sigma, E \cup B)$ into (Σ, B, \vec{E}) where \vec{E} represent the equations reoriented into rewrite rules (conditional if we have a conditional equation). Then $(\Omega, B_{\Omega}, \vec{E}_{\Omega})$ is a constructor decomposition only if the following properties are true.

1. $t =_{\Omega} t' \iff t =_B t'$
2. $t = t!_{\vec{E}_{\Omega}, B_{\Omega}} \iff t = t!_{\vec{E}, B}$
3. $C_{\Sigma_{\Omega}/E_{\Omega}, B_{\Omega}} = C_{\Sigma/E, B}$

The three statements above essentially state that there must be 1. equality on terms on both algebras, 2. the fully rewritten terms are also equivalent in the decomposition and original, and 3. the canonical algebras are equivalent.

3.6 Why Care About Decompositions?

This decomposition allows us to reason on the constructors. Doing so allows CBRL to significantly reduce the computation required to simplify goals and calculate something known as the congruence closure (discussed later). In the figure below we can see the relationship. The main thing to focus on here is the di-

$$\begin{array}{ccc}
(\Sigma, E \cup B) \longrightarrow (\Sigma, B, \vec{E}) & T_{\Sigma/E \cup B}|_{\Omega} \cong & C_{\Sigma/E, B}|_{\Omega} \\
\uparrow \scriptstyle (\Omega, E_{\Omega} \cup B_{\Omega}) \longmapsto (\Omega, B_{\Omega}, \vec{E}_{\Omega}) & \uparrow & \uparrow \\
& T_{\Omega/E_{\Omega} \cup B_{\Omega}} \cong & C_{\Omega/E_{\Omega}, B_{\Omega}}
\end{array}$$

Figure 1: Theory Inclusions (Left) and Initial Algebra Isomorphisms (Right) [11]

agonal dotted arrow on the right. In simple terms this allows us to do all the proofs using constructors. In addition to this fact often the constructor decomposition is so simple that there are no equations E_{Ω} . This further reduces the amount of computation required. The only fly in the ointment here is associativity. Associativity by itself is not decidable. Other combinations of associativity, commutativity, and identity axioms are decidable [7].

4 Constructor Based Reachability Logic

Now we shall cover the proof system, its inference rules, and how to describe invariants.

4.1 Constrained Pattern Predicates

A constrained pattern predicate is of the form $u \mid \phi$ where u is a term and ϕ is a quantifier free formula. u must be part of the canonical term algebra and thus described using constructor terms.

$$N \mid N > 0$$

In this example, this is a constrained pattern predication on a natural number N which can only apply if N is greater than 0.

4.2 Describing Invariants

To describe invariants with a pattern predicate recall our reachability formula.

$$A \rightarrow^{(*)} B$$

A and B are pattern predicates. The left side A represents the precondition and the right side B represents a midcondition. A midcondition is not the same as postcondition as it only needs to hold along the path **at some point**.

However due to the non-terminating nature of a distributed system, we are not able to actually prove an invariant. This is because all reachability formulas are vacuously true. In order to solve this problem we need to add a **stop** operator which pauses the distributed system at anytime. After this we can check our invariant.

This allows us to compute both "system invariants" and "inductive invariants". A system invariant is one that holds true eventually in the system (midcondition). An inductive invariant is one that must be true in the precondition and after taking one step (midcondition).

4.3 Comparison to Hoare Logic and Linear Temporal Logic

With this in mind one can easily see the relation between Hoare Logic and Linear Temporal Logic.

In the case of Hoare Logic we have the following.

$$\{A\}\mathcal{R}\{B\}$$

Translated into CBRL we have the following where $\mathcal{R}_{\mathcal{L}}$ represents the semantics of the programming language as a rewrite theory.

$$\{p : init \mid \phi\} \mathcal{R}_{\mathcal{L}} \{< skip : S > \mid \psi\}$$

Please refer to [11] and [1] for more information on how this is done.

In the case of LTL, we have the following equivalences.

$$A \rightarrow^{(*)} B \iff \forall Y \mid A \rightarrow \diamond B \vee \Box enabled$$

$$A \rightarrow^{(*)} [B] \iff \forall Y \mid A \rightarrow \circ B$$

$$B \rightarrow^{(*)} [B] \iff \circ B$$

4.4 T-Consistency

All reachability formulas must be something known as T-Consistent. This means that there must exist a unifier (substitution) such that the precondition is an instance of the midcondition with parameters Y under the substitution.

$$N + M \mid true \rightarrow^{(*)} O \mid true$$

In the example above N, M, O are all integers. And because we can match the sum of $N + M$ as an arbitrary integer O , this reachability formula is T-consistent.

The following example is not T-Consistent because the midcondition's term is no longer an integer but rather a boolean.

$$N + M \mid true \rightarrow^{(*)} true \mid true$$

4.5 Inference System

A proof begins and its sequents (steps, next goals) are of the form

$$[\mathcal{A}, \mathcal{C}] \vdash_T u \phi \rightarrow^{(*)} \bigvee_i v_i \mid \psi_i$$

\mathcal{A} represents the axioms you as a user input. \mathcal{C} are circularities which contain formulas that we wish to prove simultaneously. All formulas must be T-Consistent at every application of tactics otherwise the proof has failed.

4.6 Inference Rules

Below we list the inference rules used in CBRL. Further details and conditions on these inference rules can be found in [11].

1. Subsumption: Discharges trivial formulas. There are 2 cases. Either, the precondition is an instance of the midcondition with parameters Y or it is vacuously impossible.
2. Step: Use rewrite rules to evolve the system state by step. This is used internally by the auto rule.
3. Axiom: Use a trusted axiom to take "multiple rewrite steps". This is a very powerful inference rule that if used incorrectly may discharge a proof incorrectly. This is used internally by the auto rule.
4. Split: Split a constrained pattern predicate into an equivalent one
5. Case: Create new goals with a complete covering of constructor patterns
6. Substitution: Allows you to solve a conjunction of equalities in the precondition
7. Auto: Takes a step, applies all available axioms, and checks T-Consistency.

Let us give a few examples of when each inference rule is used.

Subsumption:

$$0 \mid true \rightarrow^{(*)} s(0) \mid true$$

can be subsumed by the more general pattern

$$N \mid true \rightarrow^{(*)} s(M) \mid true$$

Split:

$$N \mid true \rightarrow^{(*)} s(M) \mid true$$

can be split into the following conjunction of preconditions. While not useful in this case, split is often used to give the tool more information to proceed further with the proof. Below we split on the fact that either N is 0 or *non-zero*. This is obviously equivalent as we have the equivalence of $\phi \iff ((\phi \wedge \psi) \vee (\phi \wedge \neg\psi))$

$$(N \mid true \wedge N = 0) \vee (N \mid true \wedge N \neq 0) \rightarrow^{(*)} s(M) \mid true$$

Case:

$$N \mid true \rightarrow^{(*)} s(M) \mid true$$

can be cased upon the cover set of natural numbers. In other words all the possible constructor patterns.

$$(0 \mid true) \wedge (s(I) \mid true) \rightarrow^{(*)} s(M) \mid true$$

5 Describing Distributed Systems

To describe a distributed system for CBRL as input we use a language known as Maude. As mentioned above, Maude uses order sorted equational theories with rewrite rules to give us rewrite theories.

Below we first show a simple example of a rewrite theory and its proof in `rltool` the implementation of CBRL in Maude.

5.1 Some Important Notes

Because we are not using this Maude module normally but rather for proofs, we need to do several extra things. No builtin modules can be used. This is because the underlying implementation for things such as real numbers or integers is just the C data type. Therefore, it functions as "magic". In addition we need to disable booleans. Even if we use our own Naturals we still have the ability to check equality on them using built-ins. Therefore we need to disable these as well.

5.2 Choice: A Simple Example

This is one of the canonical examples used throughout the Maude manual. Choice is a module that turns a multiset into a singleton. Essentially pulling out non-deterministically any value in the multiset. One will notice that `fmod` (function module) only contains constructors and equations. `mod` contains rewrite rules.

```
set include BOOL off .
fmod CHOICE-DATA is
  *** Sorts
  sorts Nat MSet State Pred .
  subsorts Nat < MSet .
  *** Constructors
  op zero : -> Nat [ctor] .
  op s : Nat -> Nat [ctor] .
  op _ : MSet MSet -> MSet [ctor assoc comm] .
  op { _ } : MSet -> State [ctor] .
  op tt : -> Pred [ctor] .
  *** MSet containment
  op _=C_ : MSet MSet -> Pred [ctor] .
  vars U V : MSet . var N : Nat .
  eq U =C U = tt .
  eq U =C U V = tt .
endfm
mod CHOICE is pr CHOICE-DATA .
  vars U V : MSet .
  *** Choose a singleton
  rl [choice] : {U V} => {U} .
endm
```

5.3 Proving Choice

We would like to prove the following invariant.

$$M \mid \text{true} \rightarrow N \mid N \in M$$

In English, we can reword this. Given any multiset M our system will be able to reach a state where M has been transformed to N a natural number which is a member of multiset M . Below is the proof and its output.

```

--- Load our module and the tool
load ../../systems/choice.maude
load ../../rltool.maude

--- Select the module and methods
--- for proving, varsat stands
--- for variant satisfiability
--- not discussed here
(select CHOICE .)
(use tool varsat for unsatisfiability on
 CHOICE-DATA .)
(use tool varsat for validity on CHOICE-DATA .)
--- Declare the variables
(declare-vars (M:MSet) U (N:Nat) .)
--- Declare the terminating states/
(def-term-set ({N}) | true .)
--- Declare our invariant
(add-goal end-with-singleton :
  ({M}) | true => ({N}) | (N =C M) = (tt) .)
--- Begin the proof
(start-proof .)
--- Create new goals with a cover set
(case ({M:MSet}) on M:MSet by
  (M1:MSet M2:MSet) U (N':Nat) .)
--- Step and apply axioms
(auto .)
(auto .)
quit .

```

--- OUTPUT

```

...
Added goal(s):
  [end-with-singleton : {M:MSet} |||
true => {N:Nat} ||| tt = N:Nat =C M:MSet]
Command: add-goal end-with-singleton :
  ({M}) | true => ({N}) | (N =C M) = (tt) .

Started proof:
  [1 | {M&5:MSet} ||| true
=> {N&6:Nat} ||| tt = N&6:Nat =C M&5:MSet]
Command: start-proof .

```

Cases rule generated:

```

  [8 | {M&8:MSet M&9:MSet} |||
true => {N&6:Nat} |||
tt = N&6:Nat =C(M&8:MSet M&9:MSet)]

```

```

Action consumed 1 of 1 active
goals and generated 1 goals
Command: case({M:MSet}) on M:MSet
by(M1:MSet M2:MSet)U(N':Nat) .

```

Auto Results:

```

  [40 | {&5:MSet} |||
true => {N&7:Nat} |||
tt = N&7:Nat =C(&6:MSet &5:MSet)]
...
  [46 | {&6:MSet &8:MSet} |||
true => {N&9:Nat} |||
tt = N&9:Nat =C(&6:MSet &8:MSet &5:MSet &7:MSet)]
Action consumed 1 goals and generated 7 goals
Command: auto .

```

Proof Completed.

```

Action consumed 7 goals and generated 0 goals
Command: auto .

```

Bye.

While a full module and proof are provided here from here on full specifications and proofs will be omitted due to their length. Only the highlights will be provided

6 Simple Consensus Algorithm

Here we describe a simple consensus algorithm [6] between two parties over a channel. Each party sends a




Figure 2: Consensus Protocol

message onto an acceptor channel. The acceptor channel will then accept the first message it receives and reply to every subsequent party with the first received message. This can be seen in the image above where the woman sends message a first before the man sending b . Therefore when the acceptor channel replies to both it sends a back.

6.1 Specifying the Consensus Algorithm

We use three modules in this consensus algorithm.

1. IB00L: This is our self defined boolean operations and constructors

2. **MSGS**: This defines the messages that can be sent. Because we use two parties only we only have two constructors a and b one for each channel.
3. **STATE**: The state defines our system as a whole, as such it contains three parts, two parties and an acceptor channel.
4. **STATE-RULES**: This module contains the rewrite rules specifying the behavior of our distributed system.

Our state operator contains three parts separated by a pipe. The acceptor channel is in the middle and the parties are on the side. Each party contains two parts, the message it will send and a boolean which represents whether it has received a message from the acceptor channel.

```
op _,_|_|_,_ : Msg Bool Msg Msg Bool ->
  State [ctor metadata "8"] .
```

The rewrite rules that we have in our consensus algorithm are as follows.

1. 2 rules representing each party choosing a message
2. 2 rules representing each party proposing a message
3. 2 rules representing each party receiving a consensus message
4. A stop rule allowing us to pause the system and prove invariants

The accept rule for party a is shown below.

```
--- ~M represents message equality
cr1 [accept-a] :
  A,false | C | B,T => C,true | C | B,T
  if C ~M empty = false .
```

This conditional rule is read as follows: If party A has yet to receive a consensus message from the channel and the consensus channel is not empty, change the message of party a to C and mark it as accepted by changing *false* to *true*.

6.2 Proving the Consensus Algorithm

The system invariant we want to prove is as follows.

$$A_1, B_1 | C_1 | B_1, TT_1 \mid \text{init}(A_1, B_1 | C_1 | B_1, TT_1) \xrightarrow{(*)} A_2, B_2 | C_2 | B_2, TT_2 \mid \text{accepted?}(A_2, B_2 | C_2 | B_2, TT_2)$$

The predicate *init* checks if the system is in an initial state *empty, false | empty | empty, false* where

empty represents the empty message. *accepted?* checks if the system state has both actors accepting a consensus message and all the messages (both parties and acceptor channel) are the same.

The proof is very simple and the outline is essentially the same as the Choice proof. In fact the proof is so simple that we do not need to apply any inference rules. We only need to apply the *auto* tactic.

7 Self-Stabilizing Construction of Spanning Trees

The next algorithm was proposed by Nian-Shing Chen in 1991 [2]. The algorithm states the following. Given a connected graph, the algorithm will maintain a spanning tree with each node knowing its level in a tree. Each node is augmented with a level (depth in graph), a parent, and a list of neighboring nodes. Also the graph must contain a root node. With the information of level and a parent we are able to then build a spanning tree.

The graph may enter an unexpected state due to a perturbation in the system. In the case of this algorithm, the perturbation is represented by an incorrect level in a node. There are several cases. Firstly, a node that does not have a level of *parent* + 1 is in an error state. Secondly, a node that has level of $|V|$ (where $|V|$ is the number of vertices in our graph) is in an error state.

From an error state, the algorithm proposed uses 3 rules to return a valid state. In the figures depicting




Figure 3: Invalid Tree state [2]

the tree states, a node has a *(name, level)*. Arrows depict the spanning tree constructed by the algorithm.

7.1 Specifying the Self Stabilizing Tree Algorithm

There are two parts once again. The system description and the rules describing the recovery algorithm.




Figure 4: Valid Tree state [2]

7.1.1 Specification

A graph is represented as a list of nodes. We use ; as the concatenation operator.

```

op mtGraph : -> Graph
  [ctor metadata "28"] .
op _;_ : Graph Graph -> Graph
  [ctor assoc comm metadata "29" id: mtGraph] .
op _;_ : Graph NeGraph -> NeGraph
  [ctor assoc comm metadata "30" id: mtGraph] .
op _;_ : NeGraph Graph -> NeGraph
  [ctor assoc comm metadata "31" id: mtGraph] .

```

A node is defined as follows

```

op _||_>_::_ : iNat iNat iNat iNatList -> Node
  [ctor metadata "22"] .

```

the fields are as follows.

1. NodeID
2. Level
3. Parent
4. List of NodeIDs that are the neighbors of this node

7.1.2 Rules

There are 3 rules to the algorithm.

1. Rule 0: $L(i) \neq n \wedge L(i) \neq L(p) + 1 \wedge L(p) \neq n \rightarrow L(i) := L(p) + 1$, if a node's level is not equal to the number of the nodes in the graph and it is not equal to the level of its *parent* + 1 then we assign the level of the node to $level(parent) + 1$.
2. Rule 1: $L(i) \neq n \wedge L(p) = n \rightarrow L(i) := n$, if a node's level is not equal to the number of the nodes in the graph and its parent is equal to the number of nodes in the graph, assign the level of the node to be the number of nodes in the graph.

3. Rule 2: Let k be some neighbor of i , $L(i) = n \wedge L(k) < n - 1 \rightarrow L(i) := L(k) + 1; P(i) = k$, this is the recovery rule and it states that if a node's level is equal to the number of nodes in the graph and it has a neighbor which does not, change the parent of the node to that neighbor and its level to $level(k) + 1$.

The reasoning behind why rule two is able to recover is because we always have a root node. This root node is never modified which is why the predicate $L(k) < n - 1$ can always be triggered even if all other nodes have level n .

7.2 Proving the Self Stabilizing Tree Algorithm

The system invariant we want to prove is as follows.

$$G_1 \mid \text{connected?}(G_1) \xrightarrow{(*)} G_2 \mid \text{connected?}(G_2) \wedge \text{connected?}(\text{buildMST}(G_2))$$

It can be summarized as follows. Given a connected graph, do we reach a state after multiple rewrites do we reach a state where the new graph remains connected and after modifying the parents are we still able to build a spanning tree?

Note that unlike the consensus protocol, we do not constrain the set of initial states. This is because while we must always begin from a specific set of initial states in the consensus protocol. Here we want to prove that we are able to reach a spanning tree from **any** connected graph.

Once again, the proof is simple and of the same format as the Choice proof. We only require two autos to complete the proof.

8 Alternating Bit Protocol

Here we present the most complex, and hardest proof so far. This is also the distributed system that was able to discover the most amount of bugs in the CBRL tool (rltool).

The alternating bit protocol is a data level protocol. It models a lossy channel where a sender sends data from a sender to a receiver. The guarantees it provides are that the data will always be sent in order and will always be received (reliable data transfer).

The alternating bit protocol used here is built off of Rocha's PhD thesis on reachability analysis on rewrite theories [8].

The alternating bit protocol is defined as follows. The left side represents the sender. The right




Figure 5: ABP [8]

side represents the receiver. There are two channels in the protocol: 1. one is used to send data from sender to receiver, 2. and the other is to send an acknowledgement from the receiver to the sender.

In addition to this, both sender and receiver hold a bit value. This bit alternates between sends and receives to ensure that the receiver must receive the previously sent packet before the sender sends the next one.

Here we describe life-cycle of one message. The flip operator takes an $on \rightarrow off$ and $off \rightarrow on$

1. Sender sends: $(data_1, bit)$ on data channel
2. Receiver receives $(data_1, bit)$ and sends: (bit) as an ack
3. Sender receives ack flips its bit and sends: $(data_2, flip(bit))$
4. Receiver receives $(data_2, flip(bit))$ and sends: $(flip(bit))$
5. Repeat

8.1 Specification

The main state operator is shown below.

```

op _:_>_|_<:_:_ :
iNat Bit BitPacketQueue BitQueue Bit iNatList
-> Sys
[ctor metadata "20"] .

```

We will discuss what each entry in the **Sys** sort represents (left to right).

1. **iNat**: A natural number representing the message id we are sending
2. **Bit**: The sender bit
3. **BitPackQueue**: A queue of packets. A packet contains $(iNat, Bit)$. This packet contains the message id and the bit of the sender at the time the packet was sent.

4. **BitQueue**: A queue of bits. This queue of bits contain the acknowledgements that the receiver sends.
5. **Bit**: The receiver bit.
6. **iNatList**: A list of the received messages from the sender

Each of these function modules (**iNat**, **Bit**, and others) are quite similar and essentially only contain a concatenation operator. However equality enrichment [5] was applied to help the prover. If we had a simpler implementation of equivalence, contextual rewriting [12] would struggle more when computing the equivalence over an implication. This can be seen below.

```

--- list of naturals
fmod INAT-LIST is
pr INAT .
sort iNatList .
sort NeiNatList .
subsort iNat < NeiNatList .
subsort NeiNatList < iNatList .
op nil : -> iNatList [ctor metadata "10"] .
--- list concat
op _ _ : iNatList iNatList -> iNatList
[ctor assoc prec 61 metadata "11"] .
op _ _ : NeiNatList NeiNatList -> NeiNatList
[ctor ditto metadata "11"] .
--- list equivalence
op _~iNL_ : iNatList iNatList -> Bool
[comm metadata "12"] .
var L : iNatList .
vars P Q R S : NeiNatList .
var N M : iNat .
--- Equality enrichment
eq P ~iNL P = true .
... omitted ...
eq (s(N) P) ~iNL (s(M) Q) =
s(N) ~iNL s(M) and P ~iNL Q .
... omitted ...
eq (P 0) ~iNL (Q s(N)) = false .
--- Identity
eq L nil = L .
eq nil L = L .
endfm

```

All in all we have 16 different equality equations covering all the possible equality checks on all possible combinations of lists and natural numbers.

8.2 Rules

There are a total of ten rules that cover the different steps that a system may take. These rules are all unconditional and use matching in Maude's backend to

determine which rule will be taken at any time. We will describe these below.

1. stop: This is the stop operator for the CBRL tool
2. send-1, send-2: These rules put a new BitPacket from a sender or a Bit from a receiver onto their respective queues
3. recv-1a: Sender receives acknowledgement from sender that has the same bit .
4. recAck: This rule represents the sender receiving an acknowledgement from the receiver
5. addOutput: This rule takes a packet out of the BitPacketQueue, adds the message to the receiver's output list and flips the receiver's bit
6. recvIgnore: This rule ignores a receive of a message from the sender to the receiver
7. dropMsg: This drops a BitPacket message due to a lossy channel
8. dropAck: This drops a Bit acknowledgment due to a lossy channel
9. dup-1, dup-2: These duplicate either the BitPacket in the sender message queue or the Bit in receiver acknowledgement queue

Below we put a rule as an example.

```

--- N, N': Natural Numbers
--- B, B': Bits
--- (B,N): BitPacket (sender msg)
--- BPQ : BitPacketQueue
--- BQ : BitQueue (recv acks)
--- NL : Natural List
r1 [addOutput] :
N : B > ( B' , N' ) ; BPQ | BQ < B' : NL
=> N : B > BPQ | BQ < flip(B') : (N' NL) .

```

8.3 Invariant and Proof of the Alternating Bit Protocol

The invariant we wish to prove is an inductive invariant. Originally in Rocha's paper [8], the invariant was defined in linear temporal logic and is as follows.

$$good - queues \wedge inv \rightarrow \circ (good - queues \wedge inv)$$

Recall that in CBRL, this can be essentially rewritten as the reachability logic formula.

$$B \rightarrow^{\circ} [B]$$

inv represents the following property. The predicate *inv* is true when either

1. The list in the receiver's output is equivalent to the list all the numbers preceding the sender's message id (including the sender).
2. Otherwise the list in the receiver's output concatenated with the current sender message id is equivalent to the list of all the numbers preceding the sender's message id (including the sender) .

Here is an example. The first covers the first case. The second covers the second case. The natural list in the receiver is denoted by the list in the parentheses. Space is our concatenation operator.

```

gen-list(5) = (5 4 3 2 1 0)
gen-list(5) = 5 (4 3 2 1 0)

```

good-queues is a conjunction of 2 other predicates. The predicates used is actually dependent on the values of the sender and receiver bit. The English description will be given below.

1. If both receiver and sender bits are **on** or **off**,
a) all the bits in the receiver's acknowledgement queue must be all the same, b) the sender's message queue must either all have increasing message ids with corresponding bit flips (or is empty), or all the messages must be the same (or we have an empty queue).
2. If receiver and sender bits are different, a) the receiver's acknowledgement queue must either alternate or contain all of the same bit (or is empty), b) the sender's message queue must have all the same message bit pair (or be empty).

Throughout the proof it was impossible to just (auto .) through the proof and complete it. The lemmas used are dependent on the rule but they have the same idea behind it. For example in the *dropMsg* rule we have...

Lemma 1: If the sender message queue is empty and our invariant is true then the invariant still holds after a dropped message (which is impossible).

Lemma 2: If the sender queue is non-empty and our invariant is true, then after we drop a message we either have an empty queue and the invariant still holds true or we have a non-empty sender message queue where the invariant holds true.

In Lemma 1, we reason that the invariant must still be true because we did not change any values in the system state.

In Lemma 2, we reason that the invariant must still be true because if it was true before then we must still be true afterwards. There are two cases. If the queue has greater than one item, then the queue in question must have already satisfied the `good-queues` predicate which inductively reasons about the queue. Therefore if we drop one element it must still be true. Alternatively, if the queue has one element, then the invariant is vacuously true for an empty queue.

Below the truncated `dropMsg` proof is shown.

```

--- prove one rule at a time
(select-rls dropMsg .)
--- keyword for inductive inv
(inv good-queues-invariant to '[_:_>|_<:_:]'
on (N1 : B1 > BPQ1 | BQ1 < B1' : NL1)
  | (good-queues(N1:B1>BPQ1|BQ1<B1':NL1))
  = (true)
  /\ (inv(N1 : B1 > BPQ1 | BQ1 < B1' : NL1)
    = (true)
  .)
--- add our lemma
(add-axiom dropBPQ1 :
(N1 ... B1 > (nil).BitPacketQueue ... NL1) |
  (good-queues...) /\ (inv...) = (true) =>
([N2 ... (nil).BitPacketQueue ... NL2]) |
  (good-queues...) = (true)
  /\ (inv...) = (true) .)
(add-axiom dropBPQ2 : ...
.)
(start-proof .)
--- use our lemmas on the initial goal
((use-axioms dropBPQ1 dropBPQ2 .
dropBPQ1 dropBPQ2 on 1 .)
--- case rule on the sender message queue
(case (N:iNat ... BPQ:BitPacketQueue ...)
  on BPQ:BitPacketQueue by
    ((nil).BitPacketQueue)
    U ((B':Bit, N':iNat) ; BPQ:BitPacketQueue)
.)
(auto .)
quit .

```

In addition to these lemmas in order to simplify the proof, each rewrite rule was proved separately. This method was also applied in [10] to great effect where a whole browser specification was verified.

9 Challenges and Problems Encountered

Although CBRL has been around for several years since its inception, the implementation of the tool in Maude leaves much to be desired. The lack of QoL features such as readability, over the course of the

semester there were some bugs encountered that were critical. For brevity we will only discuss the most pathological bugs here.

During the computation of unifiers used in matching, due to the unbounded and undecidable nature of associativity we run into stack overflow error where the tool loops. This rendered proving anything impossible.

In the same vein, during the concatenation of elements into a list a similar issue occurred. We would also have a stack overflow crash. More specifically this occurred when we performed the following concatenation of non-empty lists.

```
NeList NeList -> NeList
```

There was a workaround to this by splitting our concatenation the following way

```
NeList List -> List
List NeList -> List
```

However this meant that it would be possible for a empty list constructor (`nil`) to show up in the list which is undesirable. While this does not affect the correctness of the proofs on lists it does make the proof harder.

The generalized case inference rule which uses pattern matching to split all relevant goals would also crash the tool. This is important when we have multiple goals (in the tens or even hundreds) that have the same pattern. For example we could have multiple goals like below. Where all the rules can be matched to the first one (N1)

```

(N1) | (invariant(N1)) = (true) =>
  (M1) | (invariant(M2)) = (true)
(N2 + 1) | (invariant(N2)) = (true) =>
  (M2) | (invariant(M2)) = (true)
(N3 - 1) | (invariant(N3)) = (true) =>
  (M3) | (invariant(M3)) = (true)
...

```

If we had multiple rules like these we would have had to case on each one of them individually which would make proofs very verbose and almost impossible to do if the number of active goals is 3 digits.

Lastly, we have the most important bug. Recall that a constrained pattern predicate must evaluate to true in order to be a valid goal. However, the tool failed to discharge vacuously false goals with `false = true`. It was not unusual to see goals of the form below.

```
(N1) | (true) = (false)
```

```

/\ (invariant(N1)) = (true) =>
(M1) | (invariant(M2)) = (true)

```

This made proving anything impossible. In addition to this, there were no possible workarounds. Recall that CBRL has a subsumption inference rule. While we could subsume such a goal into itself (which in itself is a bug), this goal would be added to the list of axioms that the tool could apply on the next step. In essence allowing us to prove incorrect statements that have a pattern predicate of $false = true$.

10 Comparison to Other Proof Systems

As mentioned earlier, I attempted to implement the Choice module from section 5 in two other proof systems, Isabelle and Coq. I spent around one month attempting to learn Coq and around 2 weeks learning Isabelle. In this time, I was unable to implement even the simplest system (Choice) I could fathom. While I understand this is not enough time to properly learn how to use such a complex ITP it did give critical insight for the comparison.

Isabelle and Coq are not intuitive to use in the proof of distributed systems. There is no notion of a "rewrite rule" which allows one to easily mutate a system state. In Isabelle I found an example [6] of a consensus algorithm which was proved in section 6. It uses a *stepfunction* and *execute* function to represent the changes to a system. However, unlike CBRL these functions themselves needed proofs and supporting modules to complete a specification of just the system, **not** the protocol. As for Coq, there exists old proofs of the alternating bit protocol dating back to the nineties. However even after reading the paper [4] and understanding the proof methodology detailed, understanding the Coq proof itself was beyond me.

While line count is not the most appropriate metric for comparison, I believe when we count both the proof and specification together it gives a cursory idea of how much more complex a proof would be in Isabelle and Coq.

System	CBRL	Coq	Isabelle
Consensus	<100 lines	N/A	~4000 lines
ABP	~800 lines	>4000 lines	N/A

11 Conclusion

CBRL is a very powerful and young invention. It differs from other proof systems by being theory generic and allowing the user to create very powerful rewrite theories to prove. This allows a user to very easily model system changes and steps in a distributed system via conditional and unconditional rewrite rules. However, the current implementation of the tool leaves much to be desired. From critical implementation bugs, quality of life issues, and lack of documentation, the tool currently is only usable by a very small amount of people. Namely the creator and those who have access to the creator of the tool.

However, while the problems listed above are very real, with the help from users such as myself, the tool has developed a lot over the past few months. Bugs have been ironed out and quality of life improvements have been made. The current implementation is already much more usable than before. In doing so, as we have shown here, distributed system proofs of varying complexity can be very easily done in CBRL and Maude.

References

- [1] M. Abir, M. Saxena, <https://github.com/mickyabir/cs576>.
- [2] N.-S. Chen, A self-stabilizing algorithm for constructing spanning trees.
- [3] F. Duran, J. Meseguer, Chc 3:a coherence checker tool for conditional order-sorted equational maude specifications.
- [4] E. Gimenez, An application of co-inductive types in coq: verification of the alternating bit protocol.
- [5] R. Gutierrez, C. R. Jose Meseguer, Order-sorted equality enrichments modulo axioms.
- [6] M. Kleppmann, <https://gist.github.com/ept/b6872fc541a68a321a26198b53b3896b>.
- [7] J. Meseguer, Variant-based satisfiability in initial algebras.
- [8] H. C. R. Nino, Symbolic reachability analysis for rewrite theories.
- [9] D. Plyukhin, G. Agha, Concurrent garbage collection in the actor model.
- [10] S. Skeirik, J. Meseguer, C. Rocha, Verification of ibos browser security properties in reachability logic.

- [11] A. S. Stephen Skeirik, J. Meseguer, A constructor-based reachability logic for rewrite theories.
- [12] H. Zhang, Contextual rewriting in automated reasoning.